

COCONUT: Clustering of Co-Occurring Neighboring Unambiguous Terms

Anouk Visser

Rémi de Zoeten

Cristina Garbacea

Abstract

The Abstract will be here.

1 Introduction

The introduction will be here.

2 Related Work

Recently (Mikolov et al., 2013) have shown that linguistic regularities in continuous space word representations can be identified by a vector offset method...

3 COCONUT

For learning the word representations (Mikolov et al., 2013) train an RNN with co-occurrence vectors of words. Instead of representing words by just one co-occurrence vector, we propose to train the model with multiple co-occurrence vectors for ambiguous words. The meaning of the word 'apple' can be determined by looking at its surrounding words, which could be: technology, iPhone, company for 'Apple', the company or: fruit, orchard, pie for 'apple' the fruit. COCONUT assumes that the meaning of a word is highly dependent on the words that accompany it and that the co-occurring words that define one meaning of 'apple' are more likely to co-occur with each other than two words that define two different meanings of apple ('iPhone' and 'technology' are more likely to occur together than 'iPhone' and 'orchard'). COCONUT will attempt to split the co-occurrence vector for 'apple' into two co-occurrence vectors, one containing 'iPhone', 'technology' and 'company', the other containing 'fruit', 'orchard' and 'pie'.

3.1 Co-Occurrence Vectors

We construct the co-occurrence vector for word A by computing the relatedness of word A with every other word in the vocabulary. We use the same

function for relatedness as (Guthrie et al., 1991):

$$r(x, y) = \frac{f_{xy}}{f_x + f_y - f_{xy}}$$

where f_{xy} denotes the frequency of x and y occurring together and f_x and f_y denote the frequency of x , respectively y .

3.2 Clustering

To find the two senses of a word, we apply k-means clustering to the co-occurrence vectors of the co-occurring words. COCONUT assumes that the words assigned to each cluster represent a different meaning of a word. Words that are not closely related to A do not contribute to either one of the meanings. Therefore, we will not use the co-occurrence vectors of all co-occurring words, but only those from the words that are closely related. Building a good decision process for defining when a word is closely related to another word is beyond the scope of this project and will most likely not necessarily lead to significant performance improvements. Therefore, we have decided to discard the words that have a relatedness score with A that falls in the bottom 50% of all relatedness-scores. Let the set of words that remains be called C . We can use the co-occurrence vectors of the words in C to find clusters, but these vectors will contain a lot of words that are not in C , do not occur together with A or do occur with A but not in C . We are only interested in finding clusters representing the different meanings of word A , therefore we will only use the co-occurring words in the vectors of C that are present in C .

4 Evaluation

We have evaluated the performance of COCONUT on a dataset containing X unique words, and has size X . Initially, we decided not to disambiguate the top X words, after extracting the

two senses of the words and their distance, we discarded half of the disambiguated words, leaving us with X words that were disambiguated.

4.1 Empirical Evaluation

4.2 Quantitative Evaluation

5 Conclusion

References

- Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. 2013. Linguistic regularities in continuous space word representations. Proceedings of NAACL-HLT, 746–751
- Joe A. Guthrie, Louise Guthrie, Yorick Wilks and Homa Aidinejad. 1991. Subject-dependent co-occurrence and word sense disambiguation. Proceedings of the 29th annual meeting on Association for Computational Linguistics, 146–152 Association for Computational Linguistics